

Abhishek Yadav

+91 9354180518 | yadavabhishekz@outlook.com | [linkedin/yadavabhishekz](https://www.linkedin.com/in/yadavabhishekz) | [github/yadavabhishekz](https://github.com/yadavabhishekz)

EDUCATION

Bharati Vidyapeeth (Deemed to be University) College of Engineering

Pune, Maharashtra

Bachelor of Technology in Computer Engineering

2022 – 2026

TECHNICAL SKILLS

Programming Languages: Python, C++, SQL

Frontend & Backend: React, FastAPI

AI/Machine Learning: LLM APIs(Groq, OpenAI, Gemini, Claude), Prompt Engineering, LLM Fine-Tuning, RAG, LangChain, LangGraph, Crew AI, Vector DataBase, Embeddings, Transformers. NLP, HuggingFace, NLTK, scikit-learn, TensorFlow, PyTorch, Keras

Data Analysis & Visualization: Pandas, NumPy, Scipy, Matplotlib, Seaborn, Power BI, Streamlit, Gardio

Tools & Platforms: Git, GitHub, Docker, CI/CD, GitHub Actions, AWS (EC2, S3), Antigravity, Jupyter Labs

EXPERIENCE

AI Intern

June 2025 – August 2025

Barco Electronic Systems Pvt Ltd

Noida, India

- Engineered AI-powered threat modeling tool using Python, Streamlit, and STRIDE methodology, reducing security assessment time by 70% and identifying 85%+ vulnerabilities automatically
- Orchestrated integration of 4 Large Language Model APIs (OpenAI, Claude, Gemini, Mistral) generating 50+ comprehensive threat models and attack trees weekly
- Delivered containerized solution with Docker, generating downloadable security reports in Markdown/PDF formats

PROJECTS

Budan's Brew - Coffee Shop Customer Service App | *Llama 3.1, RAG, Pinecone, Firebase, RunPod*

[Link](#)

- Engineered an AI-powered customer service system using Llama 3.1 LLM with modular agent-based architecture featuring Guard, Classification, Order Taking, and Details agents for comprehensive coffee shop operations
- Implemented Retrieval-Augmented Generation (RAG) system with Pinecone vector database to deliver accurate menu information, ingredient details, and allergen data with intelligent query processing capabilities
- Integrated market basket analysis recommendation engine using Kaggle dataset to generate personalized product suggestions and upselling opportunities integrated with order workflow
- Developed Python API backend-deployed on RunPod infrastructure, and Firebase integration supporting real-time order management and chatbot interactions

Recap — AI Meeting & Video Assistant | *FastAPI, LangChain, Groq, Gemini, ChromaDB, Docker*

[Link](#)

- Built a full-stack AI meeting assistant that transcribes YouTube/local recordings via Groq-hosted Whisper and Sarvam AI (for Hinglish), then auto-generates summaries, action items, and key decisions using LangChain LCEL and Llama 3.3-70B
- Engineered a RAG-based chat layer with ChromaDB and Gemini embeddings, isolating retrieval per meeting for accurate, hallucination-free Q&A across concurrent users
- Designed a FastAPI backend with background job processing and a single-page UI, containerized with Docker for one-command deployment

NLP-Powered Series Analysis System | *Llama 3, LoRA, DistilBERT, BART, SpaCy NER, Scrapy, Gradio*

[Link](#)

- Built end-to-end NLP pipeline processing 220+ episodes using BART zero-shot theme classification and spaCy NER-based character network analysis with NetworkX
- Fine-tuned DistilBERT for jutsu classification using custom class-weighted loss to handle 20:1 class imbalance, trained on 2,920 samples scraped via custom Scrapy pipeline
- Fine-tuned Llama-3-8B character chatbot using LoRA and 4-bit quantization, reducing memory footprint from 32GB to 6GB, and deployed full system via Gradio

EXTRACURRICULAR

Active member of the Entrepreneurship Development Cell (EDC)

- Led outreach campaigns in Corporate Relations Team, securing 2+ partnerships with 3+ industry partners and later promoted to Project Management Team for executing 4+ events with 200+ participants.

Technical Head, Association of Computer Engineering Students (ACES)

- Developed technical infrastructure and provided support for 4+ departmental events and workshops.